

GOLD-001 closed at 150

2026-08-31T03:39:01Z · closure hash 32b59f774d3efa31 · every figure read from the records at build time

READ THIS FIRST

150 cases is the achieved benchmark-size target. It is not coverage. 96 of 150 eligible cases come from one provider (openai). A per-provider result from this set is not a comparison between providers. the largest recorded category holds 58 of 150 cases and the smallest holds 1; 33 cases predate the field entirely. An unweighted score over this set is dominated by the largest category. Genuine multi-hop is still **n=1**. The preregistered calibration pilot was never run and the frozen corpus is still not reproduced, so retrieval stays blocked. None of that is fixed by reaching 150.

150
human_verified

150
holdout_eligible

9
human_rejected

1
genuine multi-hop

Where the 150 come from

group	verified	rejected	eligible	eligibility read from
001	16	2	16	evals/gold/batch_001_v2/overlay.json
002	17	1	17	evals/review/gold_review_batch_002.json
003	20	0	20	evals/review/gold_review_batch_003.json
004	14	1	14	evals/review/gold_review_batch_004_final.json
005	15	4	15	evals/review/gold_review_batch_005_final.json
006	8	1	8	evals/review/gold_review_batch_006_final.json
HA-01-HA-60	60	0	60	evals/review/gold_review_HA01_HA60_final.json
all	150	9	150	

Counts are derived through `rag_v1.gold.eligibility`, re-evaluated at build time rather than copied from the status file. Conditions checked: `human_verified`, `every_claim_has_a_deterministic_check`, `critical_strings_present_in_evidence`, `evidence_hash_valid`, `no_unresolved_scope_defect`, `required_evidence_declared`.

The 60 admitted cases

Admitted from 159 -page *Production_RAG_v1_Full_150_Case_Review.pdf*, sha256 bf6190fc53ee4ada6c948093d30e8fa7feac3dbf3300918ec75886d2a5a8f786 . Bound by **evidence identity** — version id, offsets and evidence hash — never by the short HA-nn label, because a separate 64-case packet uses the same labels for different cases. 77 spans across 60 records were re-sliced from the pinned source and rehashed before any approval was applied. The decision is the project owner's, read from `evals/review/human_decisions_HA01_HA60.json`. Codex, Grok Expert and ChatGPT each reviewed the set independently and approved nothing; an independent review is a recommendation. The separate Claude-authored 64-case packet is NOT_ADMITTED and contributed no case to this total. See GOLD-001-alternate-HA-packet-disposition.json.

HA-15 — a finding kept, not deleted

The detector reports `NONCRITICAL_ANAPHORA` : “refers to 'the model' with no antecedent in the span”. The scored fact — that the returned handoff JSON is validated locally in the SDK — does not depend on resolving that neighbouring phrase, so the case carries an explicit `project_owner` override *and keeps the finding on the record*.

HA-47 — an evidence boundary completed

Its second span opened on “It” with the antecedent stranded in the first span, which fails the independently self-contained-span rule. The two spans are replaced by the one contiguous slice that carries both, recomputed from the frozen source at 4308:4916 (608 chars):

```
e894c94d831ccfd2678f4cd132b72b52e44770d07ebeaab6c51e96e0e312a203
```

Reason `EVIDENCE_BOUNDARY_COMPLETION` , `CRITICAL_ANAPHORA_REPAIR` . The pre-repair spans and their hashes are preserved in the record's revision history and are not the admitted evidence. A paragraph break is present; read from the predicate rather than waived, no eligibility condition looks at paragraph structure, so `eligibility_blocking` is **False**.

Coverage, as measured

Provider

provider	cases	share	
openai	96	64%	
anthropic	54	36%	

Reasoning type

category	cases	share	
exact_lookup	58	39%	
(not recorded in this batch's schema)	33	22%	
configuration_interaction	23	15%	
error_behavior	19	13%	
lifecycle_compatibility_migration	8	5%	
request_response	5	3%	
lifecycle	2	1%	
ambiguity_disambiguation	1	1%	
genuine_multi_hop	1	1%	

33 eligible cases carry no `reasoning_type` at all — batches 001 and 002 predate the field. They are counted separately rather than folded into a category they were never assigned.

Evidence shape

shape	cases	share	
single_span	93	62%	
(not recorded in this batch's schema)	33	22%	
multi_span_same_fact	16	11%	
multi_span	7	5%	
multi_document	1	1%	

CONCENTRATION

The 150 cases are anchored in **60** distinct document versions, and the single most-used document supplies **14** of them (9%). A retrieval system that happens to chunk that one document well will look better than it is.

Ambiguity cases: **1**. The set cannot measure whether a system declines to answer an under-specified question.

Protocol deviation — ACCEPTED_PROTOCOL_DEVIATION

The pilot was never run. Sixty standalone drafts (HA-01 ... HA-60) were authored first, outside the preregistered sequence, and are now admitted. **These 60 cases are not the preregistered pilot and may not be described as it, retrospectively or otherwise. The pilot remains unrun.**

A prospective pilot is a commitment made before the data is seen. Its value is that the thresholds cannot be chosen to fit the result. Reviewing 60 cases afterwards, however carefully, cannot recover that property, and no amount of downstream checking converts a retrospective review into a preregistered one. What the mitigations below establish is that these particular cases are sound; what they do not establish is that the method was calibrated before it was used.

mitigation actually performed	figure	what it does not establish
structural, integrity and review-lineage checks on the derivative	1,576 / 1,576 passed	the packet states these are not proof of semantic correctness or full-corpus authenticity
deliberately corrupted negative controls	8 detected	shows the checks can fail; not a semantic guarantee
Codex semantic review	a codex_review block on every one of the 60 records	Codex authored the derivative; its review is not independent of it
Grok Expert independent semantic review	60 reviewed, 60 PASS, 0 REVISE, 0 REJECT	These are captured accessibility transcripts from the observed UI, not a signed provider API attestation.
ChatGPT independent semantic review	58 PASS, 1 pass-with-noncritical-anaphora-override (HA-15), 1 fix-required-then-approve (HA-47), 0 semantic rejections	an independent review is a recommendation; it approves nothing
re-derivation of every record from the frozen source in this session	60 records, 77 spans re-sliced at their offsets, rehashed, and version_ids recomputed through the content-derived chain	certifies these 9 document versions, not the 202-document snapshot
project-owner decision	60 approved, 0 rejected	read as an input; no script produced it
deterministic eligibility gate	60/60 holdout_eligible	metadata and evidence integrity; it does not read semantics
no retrieval leakage	retrieval_was_not_run true, systems_executed empty	none - this one is an invariant, and it holds

Corpus reproduction — still incomplete

RETRIEVAL_BLOCKED

63 OpenAI documents are individually reproducible from pinned commits. The collective snapshot `snap_689e336380a054d8039dc35b2c09cd0a` is not reproduced: **139** Anthropic documents and **2482** unbuildable identities remain outstanding.

The snapshot id is a single digest over all 202 normalized texts at once. Reproducing 63 of them certifies those 63 documents individually and says nothing about the collective fingerprint. Recovering the evidence excerpts behind any set of GOLD cases does not reproduce the corpus. Reaching 150 admitted cases does not certify the corpus those cases point into, and no retrieval system may be run until that gate clears.

Splits are not frozen

`holdout_frozen` is **false**. 150 eligible cases can support the 30–40 / 70–100 split by size, so size is no longer the reason. What is missing is a split policy: the set is provider- and category-skewed and holds 1 genuine multi-hop case, so an unweighted split would decide by accident which categories the holdout can measure. Freezing also stays blocked while corpus reproduction is incomplete.

UNTOUCHED

SYSTEM-A and SYSTEM-B remain frozen and unexecuted. `retrieval_was_not_run` is **true** and `systems_executed` is empty. No candidate selection has seen a retrieval outcome, which is the property that makes a future holdout worth having.

GOLD-001 closure at 150 · hash `32b59f774d3efa31` · generated from GOLD-001-150-case-closure.json, GOLD-001-eligibility-status.json, GOLD-001-HA-admission.json, GOLD-001-protocol-deviation-001.json and the group records. No figure in this document was typed by hand.