

CORPUS-001 — the corpus is not reproduced

2026-08-31T03:55:16Z · target `snap_689e336380a054d8039dc35b2c09cd0a` · every figure read from the records at build time

THE FINDING

139 of 202 documents have no recovered historical bytes, and the snapshot digest is not reproduced. That is the whole blocker, and it is now a single one: the 139 missing Anthropic captures. Everything else this attempt produced narrows the problem without solving it.

49 EXACT_MATCH	14 EXPECTED_HASH_UNKNOWN	139 MISSING_SOURCE	no snapshot reproduced
-------------------	-----------------------------	-----------------------	---------------------------

Two snapshot parameters recovered

Both were carried as unknown. They are confirmed *together*, by arithmetic: `experiments/EXP-007/results.json` records the manifest hash, and only one candidate name reproduces the frozen id from it.

```
stable_id("snap", "v1-openai-anthropic",
  "452479294cbbbe702b15d2df3f5a268023247dcd630782aad5e17690cee7b17",
  "v1.0", chunking_hash)
= snap_689e336380a054d8039dc35b2c09cd0a
```

Any other name gives a different 128-bit value, so the match establishes the name and the manifest hash at once. The manifest hash is the more useful of the two: it covers only the 202 `(version_id, content_hash)` pairs, so it isolates corpus *content* from the parser and chunking parameters — and it separates a content failure from a parameter failure when a candidate fails.

The ledger

status	documents	share	
MISSING_SOURCE	139	69%	
EXACT_MATCH	49	24%	
EXPECTED_HASH_UNKNOWN	14	7%	

Exact document-version reproduction **24.3%**; bytes reproduce for **31.2%**. Every `MISSING_SOURCE` row is Anthropic, and none was filled with a live page.

What survived the corpus

89 of 202 documents can be given an expected `version_id` from surviving artifacts — 40 of them Anthropic. Two oracles make those identities usable:

oracle	surviving	construction
version_id	166	stable_id("ver", src_id, content_hash)
chunk_id	803	stable_id("chk", version_id, section_path, char_start, char_end, content_hash(text))

It binds the document's version identity to its section structure, its character offsets and the chunk's own text at once, so a match cannot come from a document that merely hashes the same. It confirms 36 of the 63 reproduced OpenAI documents at that level.

THE 14 UNKNOWN OPENAI IDENTITIES

14 → **14.** the 14 are exactly the OpenAI documents that never appeared in any logged retrieval result — neither their version_id nor any of their chunk ids survives. There is no further mapping to mine; the search is exhausted by evidence, not by effort.

The host could not be searched from here

The brief asks for a sweep of Windows user directories, WSL mounts, Desktop, Downloads and Docker volumes, where the gitignored `data/raw` captures would live. This session runs in an isolated cloud VM (Claude Code remote execution) on a single root disk.

requested location	what it is	present here
<code>/mnt/c/Users</code>	Windows user directories via a WSL mount	no
<code>/mnt/d</code>	second Windows drive	no
<code>/mnt/wsl</code>	WSL interop mount	no
<code>/media</code>	removable media	present but empty — the stock mount point, not removable media
<code>/var/lib/docker</code>	Docker volumes	no
<code>/var/run/docker.sock</code>	Docker daemon	no

No Windows host, no WSL mount and no Docker daemon is reachable from this session. The gitignored `data/raw` captures were never in this container. A host-machine sweep has to be run on the host itself; its results can be brought back here as an artifact, which is how the earlier host evidence in this project arrived.

Everything reachable was swept: 6 filesystems, **0** archives or dumps found. Every provenance-anchor hit outside the repository resolves to this session's own packets, scratchpad and log.

The 40 known Anthropic identities

All 40 were located. **0** carry the document body: every one of the known identities is recorded, and not one recording carries the document body. GOLD records keep evidence spans, never the normalized text they were cut from.

WHAT THAT STILL LEAVES

16 of the 40 carry exact historical byte-slices from closed GOLD evidence — **4,551 verified bytes** at known character offsets. They cannot rebuild a document; the gaps are gone. They can refute one: a capture claiming to be the 2026-08-17 state must reproduce those bytes at those offsets before it is worth hashing.

The 139, split by what could verify them

group	documents	how a candidate is decided
A — <code>version_id</code> recorded	40	a candidate capture is accepted or rejected outright: normalize it with the frozen parser, hash it, derive <code>stable_id("ver", src_id, content_hash)</code> and require equality with the recorded identity
B — no recorded identity	99	no per-document oracle survives. A candidate can only be checked collectively: assemble all 202 (<code>version_id</code> , <code>content_hash</code>) pairs and require the manifest hash to match. That is all-or-nothing across the corpus.

candidate source	exactly verifiable	status
original local raw capture under data/raw/	yes	not present in this environment; must be searched on the host
document_version rows from the original PostgreSQL database	yes	no project database found; the local cluster holds only the three default databases
an archived project ZIP containing data/raw	yes	the one archive lead in this project resolved to an early scaffold with no corpus
a timestamped web archive capture of the 2026-08-17 state	yes	archive hosts are not reachable through this session's proxy; a capture fetched elsewhere can be verified here
provider-owned historical sources with immutable refs	yes	this is what made the OpenAI half recoverable; the Anthropic docs have no pinned form
current live platform.claude.com pages	no	NOT a recovery path. Drift was already measured at 12 of 14 sampled documents. Useful only as a comparison lead, and only after historical recovery is exhausted.

Audit correction CORPUS-001-AC-001

SUPERSEDED “the original 2,482 unbuildable identities remain a corpus reproduction blocker”
REPLACED WITH “The original 2,482 value counts failed Batch-006 authoring attempts and is not part of the corpus snapshot digest. It remains relevant only to exact authoring-pipeline reproduction.”

The 2,482 comes from one statement in `scripts/export_batch_006.py` : `removed["unbuildable"] += 1` when a builder returns `None` . They are mined spans *inside* the 202 documents, not documents. And the count is of builder attempts, not distinct spans — the generator iterates one fact list twice, and 2,482 exceeds both the 1361 facts mined and the distinct evidence texts of that same run.

The superseded wording is preserved here rather than rewritten out of the older reports. Those reports were accurate records of what the project believed when they were written; this correction is the audit trail, not a silent edit.

State

RETRIEVAL_BLOCKED CORPUS_REPRODUCTION_INCOMPLETE true, corpus_snapshot_reproduced false, RETRIEVAL_BLOCKED true. The blocker is 139 missing Anthropic document captures . GOLD is untouched at 150 verified / 150 eligible / 9 rejected / 1 genuine multi-hop. No retrieval was run, no split was frozen, no live page was fetched, and no replacement snapshot was created.
--