

GOLD-001 — Batch 004 Closure and Project Eligibility State

Production RAG v1 · closed 2026-08-23T06:17:18Z by project_owner · every candidate reached an explicit human decision

BATCH 004 IS CLOSED

14 approved, 1 rejected, 0 outstanding from 15 candidates. All 14 approved cases pass the deterministic eligibility gate. The project now holds **67 human_verified** and **67 holdout_eligible** cases across four batches, with 4 rejections preserved as audit examples.

No retrieval system has been run against any candidate in any batch. SYSTEM-A and SYSTEM-B remain frozen and unexecuted. The holdout is **not** frozen.

14 / 15
approved · acceptance 93.3%

67
project holdout-eligible
was 53 before

1 / 559
bridge pairs that became
a genuine multi-hop case

1
genuine multi-hop cases
in the whole project

1. The owner's decisions

Candidate ids are abbreviated to their final two digits; 01 is GOLD-B004-01.

id	provider	reasoning type	shape	internal review	repaired	decision
01	anthropic	error_behavior generated as configuration_interaction	single_span	NEEDS_REPAIR	text only	APPROVE
02	anthropic	configuration_interaction	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
03	anthropic	error_behavior	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
04	anthropic	error_behavior	single_span	NEEDS_REPAIR	yes	APPROVE
05	anthropic	error_behavior	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
06	anthropic	exact_lookup	single_span	NEEDS_REPAIR	yes	APPROVE
07	anthropic	lifecycle_compatibility_migration	single_span	NEEDS_REPAIR	text only	APPROVE
08	openai	ambiguity_disambiguation	multi_span	REJECT_RECOMMENDED	no	REJECT
09	openai	ambiguity_disambiguation	multi_span	NEEDS_REPAIR	yes	APPROVE
10	openai	configuration_interaction	single_span	NEEDS_REPAIR	yes	APPROVE
11	openai	error_behavior generated as configuration_interaction	single_span	NEEDS_REPAIR	text only	APPROVE
12	openai	configuration_interaction	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
13	openai	exact_lookup	single_span	NEEDS_REPAIR	text only	APPROVE
14	openai	exact_lookup	multi_span	NEEDS_REPAIR	yes	APPROVE

id	provider	reasoning type	shape	internal review	repaired	decision
15	openai	genuine_multi_hop	multi_document	NEEDS_REPAIR	yes	APPROVE

What closed. Reasoning types among the approved cases: `error_behavior` 5, `configuration_interaction` 3, `exact_lookup` 3, `lifecycle_compatibility_migration` 1, `ambiguity_disambiguation` 1, `genuine_multi_hop` 1. Evidence shapes: `single_span` 11, `multi_span` 2, `multi_document` 1. Providers: `anthropic` 7, `openai` 7.

The internal review's status and the owner's decision are different columns and were never allowed to collapse into one. Ten candidates the review marked `NEEDS_REPAIR` were approved after repair; the one it recommended for rejection was rejected on the owner's own reasoning, not on the recommendation.

REJECTED — KEPT, NOT DELETED

GOLD-B004-08 — Rejected: BENCHMARK_QUALITY / TAXONOMY_MISCLASSIFICATION. The facts are not disputed — ContentDeltaEvent carries type "content.delta" and ContentDoneEvent carries type "content.done" — but two discriminator constants are not meaningful disambiguation. Relabelling to `exact_lookup` would produce two trivial literal lookups of little benchmark value. Not salvaged to increase the candidate count. Preserved as an audit example.

2. What the repairs changed

Six candidates had an anchor repaired; four more had only their question, answer or claims rewritten. Every anchor repair is a strict outward growth of the span it replaced, both hashes are recorded, and each approval pins the post-repair hash — an approval quoting a pre-repair hash is refused by the importer.

candidate	old span	new span	reason	approved anchor
GOLD-B004-04	48081–48225	47953–48225	evidence_boundary_completion	8eb88b1c77cc0a04bd7ffe36...
GOLD-B004-06	16744–16843	16295–16843	evidence_scope_completion	b38e23b266e3c424991cafaf...
GOLD-B004-09	6938–6997, 7362–7509	6627–6997, 7072–7509	evidence_scope_completion	53b4b866ce9f1066a9bbe80d...
GOLD-B004-10	1979–2121	1827–2121	evidence_boundary_completion	352fcab03e4bb99e38aeb115...
GOLD-B004-14	— (scope span added)	1576–1654	evidence_scope_completion	6ed28e39bdd7b06085637479...
GOLD-B004-15	16756–16910	16313–16910	evidence_scope_completion	b0ef211a15b2f158ac093829...

Human overrides

A noncritical finding blocks until a person accepts it. These were accepted, and none was deleted: the detector still reports every one, and a **critical** finding cannot be overridden at all.

candidate	finding	accepted by	disposition
GOLD-B004-02	NONCRITICAL_ANAPHORA	project_owner	finding retained
GOLD-B004-05	NONCRITICAL_DEPENDENCY	project_owner	finding retained
GOLD-B004-15	NONCRITICAL_ANAPHORA	project_owner	finding retained

3. The genuine multi-hop case

GOLD-B004-15 · openai · genuine_multi_hop · multi_document · requires all evidence True · composition check **PASS**

Q. If I set `needs_approval` to `True` on a function tool used with hosted agents, what happens?

A. The tool requires approval, but SDK tool-approval interruptions are not supported there, so a function tool whose `needs_approval` is not `False` is rejected before the request is sent.

E1 · VER_AE3BFCC42C733C5051ABDA30F0F6DB07 1327–1436 (109 CHARS) · HUMAN-IN-THE-LOOP › MARKING TOOLS THAT NEED APPROVAL

Set ``needs_approval`` to ``True`` to always require approval or provide an async function that decides per call.

critical strings: `needs_approval`, `True`
72cb1e104a68797ed1296cfed17b04ba519cd1890449a2e5542f4ed36ad4e836

E2 · VER_AE909BF8B4BBBE1D1A11119447F7AC94 16313–16910 (597 CHARS) · MODELS › OPENAI MODELS › HOSTED MULTI-AGENT (EXPERIMENTAL) › LOCAL FUNCTION TOOLS

All hosted agents share the model and tools configured for the request. The Responses API decides which hosted agent calls a function. The normal SDK Runner executes the function locally and injects a ``function_call_output`` with the same call ID into the active WebSocket response, which lets the service resume the original hosted caller. Function execution still passes through the Runner's normal guardrails, hooks, and failure conversion. SDK tool approval interruptions are not supported: any function tool whose ``needs_approval`` setting is not ``False`` is rejected before the request is sent.

critical strings: `hosted agents`, `needs_approval`, `False`
b0ef211a15b2f158ac09382986237d5050567eaf25439c000f7fc76a3415bafd

The approved chain. For hosted agents, a function tool with `needs_approval` set to `True` has a value that is not `False`, so it is rejected before the request is sent rather than pausing for approval.

SCOPE, EXPLICITLY

This result holds on the hosted-agent surface and is not to be generalised beyond it. In the ordinary Runner flow, `needs_approval = True` does what span 1 says and pauses for approval. As generated, the question was unqualified and the composed answer was false in the default path; the qualification lived in a section heading, and the repair moved it inside the anchor.

4. What it cost to find one chain

The composer tested **559** bridge pairs. **1** passed the composition check; **558** were rejected.

rejection reason	pairs	share
the spans were two unrelated lookups	379	68%
no bridge relationship existed	147	26%
span 1 alone answered the full question	16	3%
span 2 alone answered the full question	16	3%
the composed answer introduced unsupported inference	0	0%

Five pairs cleared every check except the entity-state rule and were reviewed individually. All five are correct rejections, so on this evidence the rule is not too strict — it was the only check that caught them.

bridge entity	provider	verdict
OpenAIChatCompletionsModel	openai	CORRECT_REJECTION
allowed_callers	anthropic	CORRECT_REJECTION
max_tokens	anthropic	CORRECT_REJECTION
tool_result	anthropic	CORRECT_REJECTION
view_range	anthropic	CORRECT_REJECTION

Diagnostic only: 0 promoted into the batch, batch regenerated: false. Choosing candidates by re-reading the rejection list is how a benchmark ends up measuring its own generator.

5. Project eligibility

batch	candidates	human_verified	rejected	holdout_eligible	genuine multi-hop	read from
001	18	16	2	16	0	batch_001_v2
002	18	17	1	17	0	v1
003	20	20	0	20	0	v1
004	15	14	1	14	1	v1
all	71	67	4	67	1	

`human_verified` counts approvals a person gave and does not change. `holdout_eligible` counts cases a machine can still check — approval, a deterministic check for every claim, critical strings present in their evidence, valid hashes, no unresolved scope defect, and, for a multi-span case, a declaration that all of its evidence is required.

A GATE TIGHTENED MID-PROJECT REACHES BACKWARDS

The multi-span condition added for this batch, in its first form, also demanded per-span critical strings — and disqualified five closed batch-003 cases for a convention that arrived after a person had approved them. It now flags only a *mixed* record, where one span is checked and another is not. Batch 003 stands at 20. Worth stating plainly, because the failure mode is silent: the count simply comes back lower and looks like a discovery.

6. Genuine multi-hop — one observation

1 of 67 eligible cases. That is one observation. It proves the benchmark infrastructure can represent a genuine multi-hop case — anchor it, check its composition, carry it through review and eligibility — and it does **not** mean the category is adequately sampled. A single case cannot support a claim about how any system handles multi-hop reasoning.

The generation figure is the finding to carry forward: 559 bridge pairs tested, 1 valid chain. In this corpus, two facts sharing an identifier are almost never two halves of an argument, and no candidate was regenerated to improve that ratio.

7. What `precheck_holdout_ready` does and does not mean

Batch 004 produced **15 of 15** candidates precheck-ready. The source-integrity review that followed repaired **10** of them and recommended **1** for rejection.

That is not a precheck failure — the precheck is deliberately structural. It verifies hashes, offsets, string containment, anaphora and anchor size, and it means **structurally capable**: not semantic correctness, not human approval, not holdout eligibility. The review is what showed why the separation has to be maintained rather than assumed.

8. Erratum

near-miss bridge-pair count — was “3 pairs rejected only by the entity-state rule”; is **5 pairs**. The 3 came from a manual probe run mid-development, with the composer's per-run limit and used-fact set in force and before the entity-state rule existed. `scripts/diagnose_b004_near_miss.py` derives the set properly, and the PDF builder now reads the count from it rather than hardcoding one. Generation figures affected: no.

9. Provenance and invariants

what	value
reviewed-state sha256	1ce046829660332c073d41101d74415efd569ab6e2a5059c793a10d001bed712
generation batch sha256	e7c6d58c3c21b9ece940ae9665cba96ada35c887558803bba1797aaa142b93f9
closure sha256	57269018d21a2b1f72b607b419ce5f9c02789ff98e6699ca0c58ea65811c8908
corpus snapshot	snap_689e336380a054d8039dc35b2c09cd0a
validator	14 cases, 0 failures
holdout frozen	false

- The generation artifact was not rewritten. Repairs live in a separate file, and the closure records both identities, so an approval traces to the text a person saw *and* to the run it came from.
- No retrieval system was run against any candidate, in generation, review or closure. `retrieval_was_not_run = true ; systems_executed = []`.
- SYSTEM-A `9afc5b7...` and SYSTEM-B `304c3509...` remain frozen and unexecuted.
- Batches 001–003 are untouched; their closure hashes are unchanged.
- 67 eligible cases cannot support both a 30–40 validation split and a 70–100 holdout.

Generated by `scripts/build_batch_004_closure_pdf.py` from `GOLD-001-batch-004-closure.json`, `gold_review_batch_004_final.json`, `gold_review_batch_004.json`, `human_decisions_batch_004.json`, `GOLD-001-eligibility-status.json`, `BATCH-004-near-miss-multihop-review.json` and `validate_golden_batch_004.json`. Every figure is read from those artifacts at build time, and the build refuses to run if the closure hash no longer covers the records, if the closure and the batch disagree on counts, if the eligibility status disagrees with the closure, or if a holdout has been frozen. Raw provider documentation is not redistributed; quoted spans are the short excerpts under review.