

GOLD-001 — Batch 005 Closure and Project Eligibility State

Production RAG v1 · closed 2026-08-24T05:21:19Z by project_owner · every candidate reached an explicit human decision

THE FINDING: A STRUCTURAL PASS IS NOT A SEMANTIC ONE

All **19 of 19** batch-005 candidates were `precheck_holdout_ready`. The source-integrity review that followed repaired **7** of them and recommended **4** for rejection — and the owner rejected exactly those 4. The precheck verifies hashes, offsets, string containment, anaphora and anchor size. It cannot see a question whose subject is not the subject of its own evidence, and four of these had exactly that shape.

The batch also came back **short**: 30 was the target and 19 was exported. That is the second finding, and it is in this document rather than a footnote.

15 / 19

approved · acceptance 79.0%

82

project holdout-eligible
was 67 before this batch

19 / 30

exported against
the generation target

0

new multi-hop chains
from a second search

1. The owner's decisions

Candidate ids are abbreviated to their final two digits; 01 is GOLD-B005-01.

id	provider	reasoning type	shape	internal review	repaired	decision
01	anthropic	ambiguity_disambiguation	multi_document	REJECT_RECOMMENDED	no	REJECT
02	anthropic	configuration_interaction	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
03	anthropic	error_behavior	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
04	anthropic	error_behavior	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
05	anthropic	exact_lookup	single_span	NEEDS_REPAIR	text only	APPROVE
06	anthropic	lifecycle_compatibility_migration	single_span	REJECT_RECOMMENDED	no	REJECT
07	anthropic	lifecycle_compatibility_migration	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
08	anthropic	lifecycle_compatibility_migration	single_span	NEEDS_REPAIR	text only	APPROVE
09	openai	error_behavior generated as configuration_interaction	single_span	NEEDS_REPAIR	anchor	APPROVE
10	openai	configuration_interaction	single_span	REJECT_RECOMMENDED	no	REJECT
11	openai	configuration_interaction	single_span	NEEDS_REPAIR	text only	APPROVE
12	openai	configuration_interaction	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
13	openai	configuration_interaction	single_span	REJECT_RECOMMENDED	no	REJECT
14	openai	configuration_interaction	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
15	openai	configuration_interaction	single_span	NEEDS_REPAIR	text only	APPROVE

id	provider	reasoning type	shape	internal review	repaired	decision
16	openai	configuration_interaction	single_span	NEEDS_REPAIR	text only	APPROVE
17	openai	error_behavior	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE
18	openai	exact_lookup	single_span	NEEDS_REPAIR	text only	APPROVE
19	openai	lifecycle_compatibility_migration	single_span	READY_FOR_OWNER_REVIEW	no	APPROVE

The internal review's recommendation and the owner's decision are different columns and were never allowed to collapse into one. They agree on the four rejections and diverge on the seven repairs, which is what an owner reading a repaired candidate looks like — a review that the decisions matched everywhere would be indistinguishable from no review at all.

Rejected — kept, not deleted

candidate	review findings	the owner's reason
GOLD-B005-01	NOT_AMBIGUITY , NOT_A_FIELD , SCOPE_IN_HEADING	Rejected: BENCHMARK_QUALITY / TAXONOMY_MISCLASSIFICATION / SCOPE_DEFECT. <code>invalid_tool_input</code> is the same semantic concept in both sources — the tool input is invalid — and the examples differ by tool (a malformed or non-HTTP(S) URL; a malformed or over-length regex) without creating meaningful disambiguation. The spans are bare definition bullets that rely on document and section scope outside themselves. Not salvaged as two exact lookups to preserve count. Preserved as an audit example.
GOLD-B005-06	NOT_A_LIFECYCLE_STATEMENT , UNANSWERABLE_AS_ASKED	Rejected: UNANSWERABLE_AS_ASKED / CATEGORY_MISCLASSIFICATION. The evidence is a schema description of what model IDs may appear in <code>fallbacks[i].model</code> and what an empty list means. It does not answer 'Where is fallbacks supported?', and it is not a lifecycle statement merely because the phrase 'not supported' appears in a schema description. The question is not to be rewritten into a different fact under this candidate ID. Preserved as an audit example.
GOLD-B005-10	RELATION_DIRECTION , SCOPE_IN_HEADING	Rejected: RELATION_DIRECTION / UNRECOVERABLE_SCOPE. The evidence says the experimental model rejects caller-supplied <code>betas</code> overrides; it does not say <code>betas</code> overrides anything, so the generated question reverses the documented relation. Rewriting around the true subject would require identifying the experimental model, whose identity is outside the exact evidence and cannot be recovered by a minimal valid expansion. Not salvaged.
GOLD-B005-13	ANAPHORA , DUPLICATE_RELATION	Rejected: DUPLICATE_RELATION / BENCHMARK_REDUNDANCY. The fact is supported — <code>S3FilesMountPattern</code> requires broad acknowledgement because <code>mount.s3files</code> uses ambient IAM authority — but it is the adjacent sibling of GOLD-B005-12 and tests the same relation: mount pattern requires broad acknowledgement because ambient cloud authority is discovered. GOLD-B005-12 is kept; benchmark capacity is not spent on an adjacent duplicate. Preserved for audit history.

All four records remain in the batch. A rejection is evidence about the miner, and deleting it would discard the only record of what the miner got wrong — `GOLD-B005-10` in particular is the regression case for relation direction.

2. What the review changed

1 candidate had its anchor repaired; 7 of 19 had their question, answer or claims re-authored. An anchor repair is a strict outward growth of the span it replaced, both hashes are recorded, and each approval pins the post-repair hash — an approval quoting a pre-repair hash is refused by the importer.

candidate	old span	new span	reason	approved anchor
GOLD-B005-09	1609–1855	1523–1855	evidence_boundary_completion	d4a15c400fe14c77b39eb20f...

Question-authoring revisions

The miner's original wording is retained on every one of them as revision 1; nothing was overwritten.

candidate	fields	as mined	as approved
GOLD-B005-05	question	What must <code>allowed_domains</code> be?	What constraint does a request-level <code>allowed_domains</code> list have to satisfy?
GOLD-B005-08	question	Where is <code>budget_tokens</code> supported?	Is manual extended thinking with <code>budget_tokens</code> supported on Claude Sonnet 5?
GOLD-B005-09	<code>atomic_claims</code> , <code>reasoning_type</code> , <code>secondary_category</code>	—	What happens if the arguments are malformed JSON, are valid JSON but not an object (for example, <code>null</code> or a list), or contain non-standard constants such as <code>NaN</code> , <code>Infinity</code> , or <code>-Infinity</code> ?
GOLD-B005-11	question	What does <code>AWS_BEDROCK_BASE_URL</code> override?	When using the <code>bedrock(...)</code> provider, what does setting <code>AWS_BEDROCK_BASE_URL</code> change?
GOLD-B005-15	<code>answer</code> , <code>question</code>	What happens when a <code>[ComputerTool]</code> <code>[agents.tool.ComputerTool]</code> is present?	When a <code>ComputerTool</code> is present, how are <code>tool_choice="computer"</code> , <code>"computer_use"</code> and <code>"computer_use_preview"</code> treated?
GOLD-B005-16	<code>answer</code> , <code>question</code>	What happens when you use a <code>Session</code> (e.g., <code>SQLiteSession</code>)?	When you use a <code>Session</code> such as <code>SQLiteSession</code> , what usage does each call to <code>Runner.run(...)</code> return?
GOLD-B005-18	<code>answer</code> , <code>question</code>	What must <code>dispatcher</code> be?	What must an Undici-specific option like <code>dispatcher</code> be paired with?

Human overrides

A noncritical finding blocks until a person accepts it. These were accepted, and neither was deleted: the detector still reports both, and a **critical** finding cannot be overridden at all.

candidate	finding	accepted by	disposition
GOLD-B005-03	NONCRITICAL_SCOPE	project_owner	finding retained
GOLD-B005-04	NONCRITICAL_SCOPE	project_owner	finding retained

3. Target 30, exported 19

46 candidates reached the semantic self-review and 27 were dropped there, leaving 19. The drops were led by generic identifiers, claims wider than their span, and category labels the evidence did not support — none of which the structural precheck can see.

An earlier draft of this batch did reach 30 candidates. Those carried pervasive question-subject / fact-subject mismatches; the miners were corrected rather than the candidates patched, and the corrected run returned 19. Padding back to 30 would have meant keeping cases a reader could not check, so the count was the thing allowed to move.

4. The second multi-hop search found nothing new

Dependency-first: a chain may only open on a sentence that states a dependency and puts the entity in a state. 122 entities had a sentence that could open a chain; 3 dependency pairs reached the composition gates, inside a budget of 1000. 1 was a valid chain — and it is the chain batch 004 had already closed, so 0 new unique chains were exported.

bridge entity	gate	why it was rejected
input_filter	span independence	span 1 alone already answers the whole question — not multi-hop
max_tokens	semantic equivalence	max_tokens is a request parameter in span 1 and an enum value in span 2: the same string naming two different things, which is a coincidence rather than a chain

TWO SEARCHES, TWO METHODS, ONE COMPOSABLE STRUCTURE

Batch 004 tested every bridge pair — 559 of them — and found one chain. Batch 005 searched dependency-first and found the same one. That is a measured property of this frozen corpus, not a failure of either search, and it is why the project's genuine multi-hop count is 1 rather than a number a later batch can be expected to raise easily. No candidate was regenerated and no case was relabelled to fill the category.

5. What closed, by dimension

dimension	all 19 records	verified only
reasoning type	configuration_interaction 8, error_behavior 4, lifecycle_compatibility_migration 4, exact_lookup 2, ambiguity_disambiguation 1	configuration_interaction 6, error_behavior 4, lifecycle_compatibility_migration 3, exact_lookup 2
evidence shape	single_span 18, multi_document 1	single_span 15
provider	openai 11, anthropic 8	openai 9, anthropic 6

Both columns carry the labels the records hold at closure, so a category the review corrected reads as corrected here rather than as the miner first guessed; the generation report is the record of what was mined. Only the right-hand column is coverage.

Claims actually checkable: 15 of 15 verified cases carry literal critical strings, so the claim-in-evidence gate ran on every one of them. That count, not the validator's green tick, is what says whether the claims were checked.

6. Holdout eligibility

The deterministic gate was re-run at closure over the 15 verified records: **15 of 15** pass. Eligibility is derived, never asserted — every condition has to hold right now:

- `human_verified`
- `every_claim_has_a_deterministic_check`
- `critical_strings_present_in_evidence`
- `evidence_hash_valid`
- `no_unresolved_scope_defect`
- `required_evidence_declared`

The four rejected candidates fail the gate on `human_verified` and cannot become eligible by any later metadata change. That is the point of keeping approval and checkability as separate states.

7. Project eligibility

batch	candidates	human_verified	rejected	holdout_eligible	genuine multi-hop	read from
001	18	16	2	16	0	batch_001_v2
002	18	17	1	17	0	v1
003	20	20	0	20	0	v1
004	15	14	1	14	1	v1
005	19	15	4	15	0	v1
all	90	82	8	82	1	

82 eligible cases cannot support both a 30–40 validation split and a 70–100 holdout. The project is aiming at roughly 30–40 validation cases and 70–100 holdout cases; no holdout is frozen, and none should be until the count supports the split.

8. Recorded for batch 006 — not implemented

Four things batch 005 established about the generator. They are recorded as preregistration inputs rather than fixed in place: batch 005's generation artifact is a historical record, and a miner corrected retroactively would leave no evidence of what it got wrong. **No batch-006 generation was run and no miner was changed.**

defect	seen in	the check batch 006 must carry
A the bare-definition-bullet rule only inspects single-span records	GOLD-B005-01	Extend the bare-definition-bullet scope rule to every span of a multi-span record, not only records with exactly one span.
B markdown reference links survive into questions built from conditional sentences	GOLD-B005-15	Strip markdown reference links whose label is itself backticked, before a question is composed.
C prose mistaken for a section heading by the parser	GOLD-B005-11	Audit the heading parser against the corpus snapshot and record how often ordinary prose is captured as a <code>section_path</code> . Do NOT fix historical parser output in place; closed batches keep what they have.
D subject–relation direction must remain explicitly checked	GOLD-B005-10	Every generated question must be re-read against its span for relation direction: the question's subject must be the source's subject, not its object.

9. Provenance and invariants

what	value
reviewed-state sha256	c9077e4f1519a99a4b0723e3db3ed98b6f2b2a6069b8f4d84c9638ae3e99a279
generation batch sha256	37bf3509a9205637588730760dc14bf2dcaed19e03e4ed6bcd4ba7580e501af3
closure sha256	ffbf9dda40ec1554d15ade41480d0927464fb2c35a2f4a08aaa9f6fb065df4c0
corpus snapshot	snap_689e336380a054d8039dc35b2c09cd0a
validator	15 cases, 0 failures on <code>evals/gold/batch_005_projection.jsonl</code>
holdout frozen	false

- The generation artifact was not rewritten. Repairs live in a separate file and the closure records both identities, so an approval traces to the text a person saw *and* to the run it came from.
- No retrieval system was run against any candidate, in generation, review or closure. `retrieval_was_not_run = true ; systems_executed = []`.
- SYSTEM-A `9afcb5b7...` and SYSTEM-B `304c3509...` remain frozen and unexecuted.
- Batches 001–004 are untouched; their closure hashes are unchanged.
- The internal source-integrity review is Claude's own reading of the candidates. It is not independent verification, and nothing in this document treats it as such — only the owner's decisions produced `human_verified`.

Generated by `scripts/build_batch_005_closure_pdf.py` from `GOLD-001-batch-005-closure.json`, `gold_review_batch_005_final.json`, `gold_review_batch_005.json`, `human_decisions_batch_005.json`, `GOLD-001-eligibility-status.json`, `GOLD-001-batch-006-preregistration-inputs.json` and `validate_golden_batch_005.json`. Every figure is read from those artifacts at build time, and the build refuses to run if the closure hash no longer covers the records, if the closure and the batch disagree on counts, if the closure's eligibility differs from the gate re-run here, if the eligibility status disagrees with the closure, if the preregistration inputs point at a different closure, or if a holdout has been frozen. Raw provider documentation is not redistributed; quoted spans are the short excerpts under review.