

GOLD-001 — Batch 006 Closed, and the Batch 007 Contract

Production RAG v1 · closed 2026-08-25T06:50:29Z · corpus snapshot `snap_689e336380a054d8039dc35b2c09cd0a` · closure `7ff5596c755a01fc...`

BATCH 006 CLOSED · THE PROJECT STANDS AT 90

The owner approved **8 of 9** and rejected **1** (89% acceptance). All 8 approved cases passed the deterministic holdout gate — re-run at closure, not asserted.

The project now holds **90 human_verified** and **90 holdout_eligible** cases, with 9 rejected across six closed batches.

8 / 9

approved by the owner

90

project holdout-eligible

60

short of the new 150 target

699

unspent evidence spans still in the corpus

1. What the owner decided

Five were approved after revisions the owner specified, three as generated, one was rejected. **Every revision is recorded as a numbered revision attributed to the owner** — not to Claude's review — because a change someone asked for and a change the author proposed are different acts, and the record has to tell them apart.

id	provider	reasoning type	question as approved	revs
01	anthropic	exact_lookup	What credential is required to create an <code>admin</code> -role service account, and what role may a workload create?	5
02	anthropic	error_behavior	What happens if <code>role: "system"</code> is used in messages with Claude Opus 4.7?	2
03	anthropic	lifecycle_compatibility_migration	How do the newer code-execution tool versions behave on Claude Haiku 4.5?	5
04	anthropic	exact_lookup	What effort level does Claude Sonnet 5 default to on the Claude API and Claude Code?	2
05	anthropic	exact_lookup	What does <code>thinking.display</code> default to on <code>claude-mythos-5</code> and <code>claude-fable-5</code> ?	3
07	openai	configuration_interaction	What does Setting <code>audio.input.turn_detection</code> to <code>None</code> disable?	0
08	openai	lifecycle_compatibility_migration	What does the OpenAI Python SDK's temporary legacy-client compatibility path require, and how should that path be treated?	5
09	openai	exact_lookup	What does <code>RunErrorHandlerResult.include_in_history</code> default to?	0

Taxonomy corrections

Three labels were wrong and the evidence was right in every case. The anchors did not move: a taxonomy change is a change of label, and a test now enforces that.

id	was	now	why
01	configuration_interaction	exact_lookup	a requirement stated in one sentence is not two settings interacting
03	exact_lookup	lifecycle_compatibility_migration	accepts / unavailable / behaves-like is a compatibility statement, not a lookup
08	configuration_interaction	lifecycle_compatibility_migration	a compatibility path described as a migration bridge is lifecycle, not interaction

Two of these were compound facts stated in a single span. The owner was explicit that neither is multi-hop, and none was labelled as such.

The rejection

GOLD-B006-06 — DUPLICATE FACT / BENCHMARK REDUNDANCY

What does `AWS_BEDROCK_BASE_URL` override?

This uses the regional ``https://bedrock-mantle.<region>.api.aws/openai/v1`` endpoint. The region can also come from ``AWS_REGION`` or ``AWS_DEFAULT_REGION``, and ``AWS_BEDROCK_BASE_URL`` can override the endpoint.

The fact is supported. `GOLD-B005-11` , approved in batch 005, already carries the same operational relation from the OpenAI *Python* library; this obtains it from the *TypeScript/JavaScript* library. Useful source corroboration, not a second benchmark case.

Duplicate control could not see it. It compares question text, span offsets and span text, and two libraries documenting the same behaviour share none of those. Recorded as defect E for batch 007 rather than patched here. The record is preserved as a negative audit example and was not replaced.

2. Why the batch was small — and why that is not a corpus problem

Batch 006 was commissioned at **28** and exported **9**. Of the 50 candidates that reached the semantic self-review, 40 were dropped. Nothing was padded in.

the census behind the shortfall	count
facts mined	1361
distinct evidence spans the miners reach	773
of those, unspent by any closed batch	699

THE FINDING THAT SETS THE NEXT BATCH'S METHOD

The corpus is not exhausted. The authoring is. 699 distinct spans in the frozen snapshot have never been used, and no deterministic template could turn them into a question without paraphrasing — what remains is long, multi-clause prose, and a template that fits it is a template that invents wording.

Refusing all paraphrase is precisely what keeps those facts out of the benchmark. That is the problem batch 007 is preregistered to solve.

3. Heading parser audit

44 of 5857 parsed headings (0.75%) read as prose rather than as a label, across 15 of 202 documents; 82 were suspicious on at least one rule. 44 of 5857 parsed headings (0.75%), in 15 of 202 documents. That is isolated. It does not justify a parser experiment on its own, and the batch-006 rule — never trust `section_path` for scope — is the cheaper fix.

Nothing was rewritten. No heading changed, no document was reparsed, no anchor moved. What changed is a rule: `section_path` is not trusted for claim scope, and a candidate's exact evidence must carry the scope its claim needs.

4. Multi-hop

No search was run. §7: batch 004 tested 559 bridge pairs and found 1 chain; batch 005 searched dependency-first and found the same one. The corpus has been measured twice. No search was run here, and no multi-span case was relabelled to raise the count. Genuine multi-hop cases in this batch: **0**. No multi-span case was relabelled to raise the count.

5. Project state

batch	candidates	verified	rejected	holdout-eligible
001	18	16	2	16
002	18	17	1	17
003	20	20	0	20
004	15	14	1	14
005	19	15	4	15
006	9	8	1	8
all	99	90	9	90

Batch 007 — preregistered, not generated

PREREGISTERED — no batch-007 candidate has been generated

THE LINE

The authoring model may change WORDING. It may not change MEANING. It authors the QUESTION; it never invents the FACT.

Batch 007 introduces **controlled evidence-grounded question paraphrasing**: where a deterministic template cannot express an explicit evidence fact, a model may author the question. This is an **authoring** change. The evidence stays frozen and exact, the ground truth is still read out of the source, and every existing gate still runs.

The order is the safeguard

- 1. frozen source evidence selected
- 2. literal source fact extracted
- 3. subject / relation / object recorded
- 4. atomic claims anchored
- 5. only then the natural question is paraphrased

Never: invent question → search for supporting evidence. Inventing a question and then hunting for evidence to support it is how a benchmark ends up testing what its author imagined rather than what the documentation says.

The entailment self-check — any failure drops the candidate

check	it fails when
A Does the exact evidence support the complete answer?	any part of the answer is not derivable from the anchored spans
B Does every atomic claim map to literal evidence?	a claim has no span whose text supports it
C Did the question introduce any new condition?	the question adds an 'if', 'when' or 'unless' the source lacks
D Did it broaden model/provider/platform scope?	the source names a model, provider, platform or surface and the question drops or generalises it
E Did it reverse the relation?	the question's subject is the source's object and vice versa, and the relation is not symmetric
F Did it introduce a causal claim absent from the source?	the answer explains why, and the source only says what
G Could the answer be verified using only the exact evidence?	verifying it needs the surrounding document, a heading, or outside knowledge

There is no flag-and-continue branch, deliberately. Every paraphrased candidate also carries its `source_fact_literal` and its subject/relation/object beside the authored question, so a reviewer can see the gap between fact and question and disagree with it.

Answer conservatism

The question may be naturalised more aggressively than the answer. Prefer answers close to the source wording. Source says *A overrides B* → answer *A overrides B*, **not** *A overrides B because this improves reliability*. The reason is not in the source. An answer that explains is an answer that asserts something the evidence cannot check.

A pilot gates the lane

10 spans that failed batch 006 *only* because no builder could express them. Run through the new paraphraser and every semantic check, then **independently reviewed** before the lane may scale.

critterion	threshold
independently judged factually sound	≥ 8 of 10
unsupported claims	0
relation-direction reversals	0
scope broadening	0
wording cleanup needed	acceptable

If it fails: Do not scale the paraphrasing lane. Revise the authoring contract first and re-pilot.

Every existing gate still runs

gate	implemented in	behaviour
bare definition scope	<code>rag_v1.gold.scoping</code>	every span, independently
critical anaphora	<code>rag_v1.gold.anaphora</code>	blocks; noncritical flags
subject / relation direction	<code>rag_v1.gold.relations</code>	REVERSED and SUBJECT_MISMATCH both drop
question form matches evidence form	<code>rag_v1.gold.questionform</code>	negative-as-positive and truncated predicates drop
duplicate detection	<code>scripts/export_batch_007.py</code>	question text, span offsets, span text — and now the relation triple
critical strings	<code>rag_v1.gold.normalisation</code>	every one must be literally inside its own span
evidence hashes	<code>scripts/export_batch_007.py</code>	each span hashes to its own text
scope self-containment	<code>rag_v1.gold.scoping</code>	section_path is never claim scope
example-code restriction	<code>scripts/export_batch_007.py</code>	a sample configuration is not a rule
evidence size	<code>scripts/export_batch_007.py</code>	<500 preferred, 1000 soft cap, 1500 hard cap
provider / model scope	<code>scripts/export_batch_007.py</code>	a scoped source needs a scoped question
holdout eligibility	<code>rag_v1.gold.eligibility</code>	deterministic, run after owner approval

Defects to fix before batch 007 authors anything

defect	seen in	proposed fix
E cross-library duplicate facts are invisible to duplicate control	<code>GOLD-B006-06</code>	Compare candidates on their (subject, relation, object) triple, normalised, in addition to text and offsets. Batch 006 records that triple on every candidate, so the material for the check now exists. Flag rather than auto-drop: two libraries genuinely differing in behaviour is a real case, and only a reviewer can tell the two apart.

defect		seen in	proposed fix
F	compound	GOLD-B006-01, GOLD-B006-03, GOLD-B006-08	Classify from the whole sentence rather than from the matched verb: a span naming a support status, a version or a migration is a lifecycle case whatever its verb, and <code>configuration_interaction</code> should require two settings that bear on each other rather than one requirement with two identifiers in it.
	single-span facts are labelled by their first verb		
G	questions inherit the breadth of their frame, not of their evidence	GOLD-B006-02, GOLD-B006-04, GOLD-B006-05	When the source sentence carries a model, platform or surface qualifier, the question must carry it too and it must appear in the critical strings. Add a generation gate: a question whose evidence is scoped and whose wording is not, does not export.

The arithmetic, which is not a plan

The project holds **90** and the target is now **150**, leaving **60**. Batch 007 targets **35-40** candidates, which at the observed acceptance rates lands between **117** and **130** — **short of 150**, so more than one batch will be needed.

A projection, not a plan. It is recorded so nobody has to compute it during review, and it must not influence any individual approval. If controlled paraphrasing yields only 25 strong cases, batch 007 returns 25.

Invariants

- No retrieval system was run against any candidate. `retrieval_was_not_run = true`; `systems_executed = []`.
- SYSTEM-A `9afcb5b7...` and SYSTEM-B `304c3509...` remain frozen and unexecuted.
- The corpus snapshot, chunks, embeddings and retrieval architecture are unchanged.
- Batches 001–006 are closed and hash-covered; the batch-006 generation artifact was not rewritten — repairs live beside it.
- No holdout and no validation split is frozen.
- Only the project owner. No AI may set `human_verified`, and controlled paraphrasing does not make Claude authoritative about anything.

Generated by `scripts/build_batch_006_closure_pdf.py` from the batch-006 closure, its composed reviewed-state file, its generation artifact, the batch-007 preregistration and the project-wide eligibility status. Every figure is read from those artifacts at build time. The build refuses to run if the closure's totals disagree with the records, if the eligibility gate re-run disagrees with the closure, if a taxonomy change moved an anchor, if any approved record fails the scope or relation-direction checks, if a batch-007 candidate exists, or if the page would claim batch 007 reaches the project target. Raw provider documentation is not redistributed; quoted spans are the short excerpts under review.