

GOLD-001 — Batch 006

The Coverage Push That Came Back Short

Production RAG v1 · generated 2026-08-24T06:01:08Z · corpus snapshot `snap_689e336380a054d8039dc35b2c09cd0a` · nothing here is gold

THE FINDING: 9 OF 28, AND 100 IS OUT OF REACH

Batch 006 was commissioned as the final coverage push — 28 candidates, to carry the project from **82** confirmed holdout-eligible cases past **100**. It exported **9**.

Even if an independent review and the owner approved *every one*, the project would reach **91** — still **9 short**. At the 79% acceptance rate of the weakest closed batch it reaches 89. **This batch does not get GOLD-001 to its minimum**, and no approval decision should be taken as though it might.

Nothing was padded to close the gap. Every one of the 9 cleared the structural precheck and the semantic self-review; 40 more were dropped and are listed below with the gate that caught each.

9 / 28

exported against target

91 / 100

best case if all were approved

40

dropped by the self-review

4 / 4

preregistered fixes implemented first

1. Why it came back short

Not the corpus. A census of the frozen snapshot found **699 distinct evidence spans that no closed batch has spent**, out of 773 distinct spans the miners reach. The shortage is in *authoring*: of 1364 mined facts, **2482** reached no builder that could turn them into a question without paraphrasing. The material in this corpus is long, multi-clause prose, and a template that fits it is a template that invents wording.

A new predicate lane was added for this batch — one question frame per verb the corpus actually uses, with the question built from the sentence's own subject — and it is what produced most of the 9. Pushing further meant either loosening a gate that exists because a real defect got through a previous batch, or inventing question templates that paraphrase. Neither was done.

where candidates were lost	count
unbuildable	2482
dropped by semantic review	40
duplicate evidence	36
duplicate question	26
short of target before overflow	19
duplicate evidence text	5
excluded known failure case	1
blocking anaphora	1
document concentration	1
not selected diversity	1

2. The four preregistered fixes, implemented before anything was authored

fix	from	implemented in
A bare-definition-bullet scope applied to every span	GOLD-B005-01	<code>rag_v1.gold.scoping</code>
B markdown links stripped from questions and answers	GOLD-B005-15	<code>rag_v1.gold.normalisation.strip_markdown_links</code>
C heading parser audited; section_path not trusted for scope	GOLD-B005-11	<code>scripts/audit_heading_parser.py</code>
D source and question triples recorded and compared	GOLD-B005-10	<code>rag_v1.gold.relations</code>

Each was recorded in batch 005's closure as a preregistration input, and each has a regression test built from the candidate that motivated it. Batch 005's artifacts are unchanged: the fixes are forward-looking, which is the point of recording them rather than patching a closed batch.

WHAT THE FIXES ACTUALLY CAUGHT IN THIS BATCH

A gate that never fires is indistinguishable from a gate that was never wired in, so the counts are reported rather than the intention. **12** candidates were dropped for bare definition bullets (Fix A — the rule batch 005 applied only to single-span records), **9** because they could not state their own subject and relation (Fix D), and **1** for asking about the wrong subject.

gate that fired	candidates
BARE_DEFINITION_SCOPE	12
NO_TRIPLE	9
GENERIC_IDENTIFIER	9
CATEGORY	5
QUESTION_SCOPE	2
QUESTION_FORM	1
SUBJECT_MISMATCH	1
CLAIM_SCOPE	1

3. What was dropped

gate	reasoning type	question
QUESTION_FORM	exact lookup	What must <code>openssl</code> be?
BARE_DEFINITION_SCOPE	exact lookup	What does the <code>is_enabled</code> option control?
NO_TRIPLE	exact lookup	What is the <code>nest_handoff_history</code> option?
CATEGORY	configuration interaction	What does <code>created_at</code> require?
NO_TRIPLE	configuration interaction	What happens if recent raw memories exceed <code>max_raw_memories_for_consolidation...</code>
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>default_message</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>encrypted_index</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>request_too_large</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>parsed</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>arguments</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>arguments_delta</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>refusal</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>refusal</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>group_id</code> option?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>insert_line</code> option?
NO_TRIPLE	lifecycle compatibility migration	What should I move to instead of <code>httpx</code> ?
BARE_DEFINITION_SCOPE	exact lookup	What is the <code>tool_use_id</code> option?
SUBJECT_MISMATCH	configuration interaction	What happens when the memory tool is present in your request's <code>tools</code> ?
GENERIC_IDENTIFIER	exact lookup	What must <code>name</code> be?
GENERIC_IDENTIFIER	exact lookup	What must <code>name</code> be?
GENERIC_IDENTIFIER	exact lookup	What must <code>name</code> be?
QUESTION_SCOPE	configuration interaction	What does <code>RunConfig.session_settings</code> override?

40 drops in total; the first 22 are shown. They are recorded rather than regenerated away — what the miner gets wrong is part of what this batch measures.

4. The 9 candidates

id	provider	reasoning type	question	chars
01	anthropic	<code>configuration_interaction</code>	What does Creating an <code>admin</code> -role service account require?	181
02	anthropic	<code>error_behavior</code>	What does Claude Opus 4.7 reject?	72
03	anthropic	<code>exact_lookup</code>	What does Claude Haiku 4.5 accept?	345
04	anthropic	<code>exact_lookup</code>	What does Claude Sonnet 5 default to?	76
05	anthropic	<code>exact_lookup</code>	What does <code>thinking.display</code> default to?	182
06	openai	<code>configuration_interaction</code>	What does <code>AWS_BEDROCK_BASE_URL</code> override?	206
07	openai	<code>configuration_interaction</code>	What does Setting <code>audio.input.turn_detection</code> to <code>None</code> disable?	81
08	openai	<code>configuration_interaction</code>	What does the OpenAI Python SDK's temporary legacy-client compatibility path require?	324
09	openai	<code>exact_lookup</code>	What does <code>RunErrorHandlerResult.include_in_history</code> default to?	62

Provider {'anthropic': 5, 'openai': 4} · 8 distinct documents · evidence mean 169.9, median 181, max 345 characters, 0 over the 1000-character soft cap. 9 of 9 are `precheck_holdout_ready`, which means **structurally checkable and nothing more**.

Subject and relation, on every record

Fix D's contract: a candidate states what its source says and what its question asks, and export compares them. A question that reverses its source cannot leave.

id	source subject	relation	source object	question subject
01	Creating an <code>admin</code> -role service account	<code>requires</code>	an interactive credential (a user OAuth token or a Conso	Creating an <code>admin</code> -role service account
02	Claude Opus 4.7	<code>rejects</code>	<code>role: "system"</code> in <code>messages</code> with a 400 error.	Claude Opus 4.7
03	Claude Haiku 4.5	<code>accepts</code>	the <code>code_execution_20260120</code> and <code>code_execution_202605</code>	Claude Haiku 4.5
04	Claude Sonnet 5	<code>defaults_to</code>	high effort on the Claude API and Claude Code.	Claude Sonnet 5
05	<code>thinking.display</code>	<code>defaults_to</code>	"omitted" on <code>claude-mythos-5</code> and <code>claude-fable-5</code> , t	<code>thinking.display</code>
06	<code>AWS_BEDROCK_BASE_URL</code>	<code>overrides</code>	the endpoint.	<code>AWS_BEDROCK_BASE_URL</code>
07	Setting <code>audio.input.turn_detection</code> to <code>None</code>	<code>disables</code>	automatic turn detection.	Setting <code>audio.input.turn_detection</code> to
08	The OpenAI Python SDK's temporary legacy-client compatib	<code>requires</code>	an explicit <code>httpx</code> installation and should be treated a	The OpenAI Python SDK's temporary legacy
09	<code>RunErrorHandlerResult.include_in_history</code>	<code>defaults_to</code>	<code>True</code> .	<code>RunErrorHandlerResult.include_in_histor</code>

5. Every candidate in full

GOLD-B006-01 · anthropic · configuration_interaction

Q. What does Creating an `admin`-role service account require?

A. Creating an `admin`-role service account requires an interactive credential (a user OAuth token or a Console session) — a workload may only create `developer`-role service accounts.

Admin · `ver_c299b58fe1f5...` · 441865–442046 · 181 chars · triple read by named relation

```
Creating an `admin`-role service account requires
an interactive credential (a user OAuth token or a Console session) — a
workload may only create `developer`-role service accounts.
```

critical strings: `admin`, `developer`

GOLD-B006-02 · anthropic · error_behavior

Q. What does Claude Opus 4.7 reject?

A. Claude Opus 4.7 rejects `role: "system"` in `messages` with a 400 error.

Migration guide · `ver_a7bda3595f2c...` · 55306–55378 · 72 chars · triple read by named relation

```
Claude Opus 4.7 rejects `role: "system"` in `messages` with a 400 error.
```

critical strings: `messages`

GOLD-B006-03 · anthropic · exact_lookup

Q. What does Claude Haiku 4.5 accept?

A. Claude Haiku 4.5 accepts the `code_execution_20260120` and `code_execution_20260521` tool types, but programmatic tool calling and the REPL state persistence that depends on it aren't available on it, so the newer versions behave like `code_execution_20250825` there.

Code execution tool · `ver_f65938c74d40...` · 3971–4316 · 345 chars · triple read by named relation

```
Claude Haiku 4.5 accepts the `code_execution_20260120` and `code_execution_20260521` tool types, but programmatic tool calling and the
REPL state persistence that depends on it aren't available on it, so the newer versions behave like `code_execution_20250825` there.
* `code_execution_20260521` is the same runtime as `code_execution_20260120`.
```

critical strings: `code_execution_20260120`, `code_execution_20260521`, `code_execution_20250825`

GOLD-B006-04 · anthropic · exact_lookup

Q. What does Claude Sonnet 5 default to?

A. Claude Sonnet 5 defaults to `high` effort on the Claude API and Claude Code.

Effort · `ver_238961530ce6...` · 5123–5199 · 76 chars · triple read by named relation

```
Claude Sonnet 5 defaults to `high` effort on the Claude API and Claude Code.
```

critical strings: `high`

GOLD-B006-05 · anthropic · exact_lookup

Q. What does `thinking.display` default to?

A. `thinking.display` defaults to `"omitted"` on `claude-mythos-5` and `claude-fable-5`, the same as on Claude Mythos Preview; set `display: "summarized"` to receive readable summaries.

Migration guide · `ver_a7bda3595f2c...` · 15155–15337 · 182 chars · triple read by named relation

```
`thinking.display` defaults to `omitted` on `claude-mythos-5` and `claude-fable-5`, the same as on Claude Mythos Preview; set `display:
"summarized"` to receive readable summaries.
```

critical strings: `thinking.display`

GOLD-B006-06 · openai · configuration_interaction

Q. What does `AWS_BEDROCK_BASE_URL` override?

A. This uses the regional `https://bedrock-mantle.<region>.api.aws/openai/v1` endpoint. The region can also come from `AWS_REGION` or `AWS_DEFAULT_REGION`, and `AWS_BEDROCK_BASE_URL` can override the endpoint.

OpenAI TypeScript and JavaScript API Library · `ver_f30a6447e4df...` · 17117–17323 · 206 chars · triple read by named relation

```
This uses the regional `https://bedrock-mantle.<region>.api.aws/openai/v1` endpoint. The region can also come from `AWS_REGION` or
```

`AWS\_DEFAULT\_REGION`, and `AWS\_BEDROCK\_BASE\_URL` can override the endpoint.

critical strings: `AWS_BEDROCK_BASE_URL`, `AWS_REGION`, `AWS_DEFAULT_REGION`

GOLD-B006-07 · openai · configuration_interaction

Q. What does Setting `audio.input.turn_detection` to `None` disable?

A. Setting `audio.input.turn_detection` to `None` disables automatic turn detection.

Realtime agents guide · `ver_14a2187cf221...` · 7318–7399 · 81 chars · triple read by named relation

Setting ``audio.input.turn_detection`` to ``None`` disables automatic turn detection.

critical strings: `audio.input.turn_detection`, `None`

GOLD-B006-08 · openai · configuration_interaction

Q. What does the OpenAI Python SDK's temporary legacy-client compatibility path require?

A. The OpenAI Python SDK's temporary legacy-client compatibility path requires an explicit `httpx` installation and should be treated as a migration bridge.

Release process/changelog · `ver_de67d790db97...` · 1996–2320 · 324 chars · triple read by named relation

The OpenAI Python SDK's temporary legacy-client compatibility path requires an explicit ``httpx`` installation and should be treated as a migration bridge.

- Local MCP HTTP customization continues to follow the installed MCP package: MCP Python SDK v1 supplies and uses legacy ``httpx``, while MCP Python SDK v2 uses ``httpx2``.

critical strings: `httpx`, `httpx2`

GOLD-B006-09 · openai · exact_lookup

Q. What does `RunErrorHandlerResult.include_in_history` default to?

A. `RunErrorHandlerResult.include_in_history` defaults to `True`.

Running agents · `ver_2c60e99cfd92...` · 29933–29995 · 62 chars · triple read by named relation

``RunErrorHandlerResult.include_in_history`` defaults to ``True``.

critical strings: `RunErrorHandlerResult.include_in_history`, `True`

6. Heading parser audit (Fix C)

44 of 5857 parsed headings (0.75%) read as prose rather than as a label, across 15 of 202 documents — GOLD-B005-11 's "configured through AWS_REGION, AWS_DEFAULT_REGION, or your AWS profile." among them.

44 of 5857 parsed headings (0.75%), in 15 of 202 documents. That is isolated. It does not justify a parser experiment on its own, and the batch-006 rule — never trust section_path for scope — is the cheaper fix.

Nothing was rewritten. No heading was rewritten and no document was reparsed into storage. No existing evidence anchor moved; closed batches are unchanged. What changed is a rule, not a record: section_path is not trusted for claim scope. A candidate's exact evidence must contain the scope its claim needs, and a candidate whose scope would depend on a suspicious heading is repaired or dropped.

7. Multi-hop: no search was run

§7: batch 004 tested 559 bridge pairs and found 1 chain; batch 005 searched dependency-first and found the same one. The corpus has been measured twice. No search was run here, and no multi-span case was relabelled to raise the count. Exported chains: 0. The project's genuine multi-hop count stays at 1.

8. Project state — batch 006 counts for nothing yet

batch	candidates	human_verified	rejected	holdout_eligible
001	18	16	2	16
002	18	17	1	17
003	20	20	0	20
004	15	14	1	14
005	19	15	4	15
confirmed	90	82	8	82
006 (unverified)	9	0	0	0

No batch-006 candidate is human_verified , and none is counted as confirmed anywhere in this project's records. The next step is independent review.

9. Invariants

- No retrieval system was run against any candidate at any point. retrieval_was_not_run = true ; systems_executed = [] . No question was selected, ordered or worded because of retrieval difficulty.
- SYSTEM-A 9afcb5b7... and SYSTEM-B 304c3509... were verified frozen before generation began and were not executed.
- The corpus snapshot, chunks, embeddings and retrieval architecture are unchanged. This was evaluation-authoring work only.
- Batches 001–005 are untouched; their closure hashes still cover their records.
- No holdout and no validation split is frozen.
- The semantic self-review is Claude reading its own output. It is not independent verification, and nothing in this document treats it as such.

Generated by scripts/build_batch_006_pdf.py from gold_review_batch_006.json, GOLD-001-batch-006-generation-report.json, GOLD-001-heading-parser-audit.json and GOLD-001-eligibility-status.json. Every figure is read from those artifacts at build time. The build refuses to run if any record claims a verification it does not have, if the batch records a retrieval run, if a multi-hop search ran, if any exported record fails the scope, link or relation-direction checks, if the report disagrees with the batch, or if the page would claim the project reaches 100 when the arithmetic says otherwise. Raw provider documentation is not redistributed; quoted spans are the short excerpts under review.