

# GOLD-001 — Batch 007 Fixes E/F/G, and the Pilot That Could Not Run

Production RAG v1 · 2026-08-30T03:58:20Z · corpus snapshot `snap_689e336380a054d8039dc35b2c09cd0a` · commit `c65d5204901e`

## THE FINDING

**The calibration pilot did not run.** The three preregistered generator defects are implemented and verified against the real candidates that revealed them. The pilot the preregistration requires before the paraphrasing lane may scale could not be run: the frozen evidence it must draw from is not present in this environment, and no substitute for it is admissible. **The paraphrasing lane does not scale, and no batch-007 candidate has been authored.**

3 / 3

preregistered fixes implemented

9 / 9

whole-sentence labels agreeing with the owner

0 / 10

pilot cases authored — input unobtainable

90

project holdout-eligible, unchanged

## 1. State verified before anything was written

| reads                                 | value                       |                   |
|---------------------------------------|-----------------------------|-------------------|
| human_verified                        | 90                          | PASS              |
| holdout_eligible                      | 90                          | PASS              |
| human_rejected                        | 9                           | PASS              |
| genuine multi-hop                     | 1                           | PASS              |
| closure sha256                        | 7ff5596c755a01fc57cf59b0... | recomputed, PASS  |
| retrieval_was_not_run                 | true                        | PASS              |
| systems_executed                      | []                          | PASS              |
| SYSTEM-A / SYSTEM-B                   | 9afcb5b7... / 304c3509...   | frozen, unchanged |
| batch-007 candidate or pilot artifact | none                        | PASS              |

The closure hash was **recomputed** from the nine closed batch-006 records rather than read off the closure, so the state is checked rather than quoted.

## 2. E — cross-library duplicate facts

`src/rag_v1/gold/factidentity.py`

compares the normalised (subject, relation, object) triple in addition to text and offsets; flags, never drops; reports 'not\_comparable' when no triple can be read

**The pair the owner caught is now visible.** `GOLD-B005-11` (OpenAI Python library) and `GOLD-B006-06` (TypeScript/JavaScript library) share no question text, no span offsets and no span text — which is exactly why the old comparison could not see them. Both now normalise to the triple `('aws_bedrock_base_url', 'override', 'endpoint')`, and the second is flagged.

```
B005-11 Pass `base_url` to `bedrock(...)` or set `AWS_BEDROCK_BASE_URL` to override the derived `https://bedrock-mantle.<region>.api.aws/openai/v1` endpoint.
B006-06 ... and `AWS_BEDROCK_BASE_URL` can override the endpoint.
```

Batch 005 predates the triple fields, so its triple is **derived from its frozen evidence** — never from its question, because the evidence is the part that cannot drift. Within batch 006 the check raises **0** false positives. It **flags and never drops**: two libraries genuinely differing in behaviour is a real case, and only a reviewer can tell that apart from a restatement.

### 3. F — reasoning type read from the whole sentence

src/rag\_v1/gold/reasoningtype.py

classifies from the whole sentence; lifecycle first whatever the verb; configuration\_interaction requires an interaction verb, so one requirement naming two identifiers is a lookup

| id | generated                 | owner                             | whole-sentence                    |                    |
|----|---------------------------|-----------------------------------|-----------------------------------|--------------------|
| 01 | configuration_interaction | exact_lookup                      | exact_lookup                      | PASS<br>relabelled |
| 02 | error_behavior            | error_behavior                    | error_behavior                    | PASS               |
| 03 | exact_lookup              | lifecycle_compatibility_migration | lifecycle_compatibility_migration | PASS<br>relabelled |
| 04 | exact_lookup              | exact_lookup                      | exact_lookup                      | PASS               |
| 05 | exact_lookup              | exact_lookup                      | exact_lookup                      | PASS               |
| 06 | configuration_interaction | configuration_interaction         | configuration_interaction         | PASS               |
| 07 | configuration_interaction | configuration_interaction         | configuration_interaction         | PASS               |
| 08 | configuration_interaction | lifecycle_compatibility_migration | lifecycle_compatibility_migration | PASS<br>relabelled |
| 09 | exact_lookup              | exact_lookup                      | exact_lookup                      | PASS               |

9/9 agreement with the owner's labels, including all 3 the owner had to relabel. Lifecycle is tested first, whatever the verb: a compatibility sentence almost always also contains a lookup verb, and reading it as a lookup is precisely how GOLD-B006-03 was mislabelled.

## 4. G — a scoped source needs a scoped question

src/rag\_v1/gold/questionscope.py

a scope qualifier in the evidence must appear in the question AND in the critical strings, else the candidate does not export; a comparative aside is not a scope

| id | as generated                                         | as owner-approved       |
|----|------------------------------------------------------|-------------------------|
| 02 | SCOPE_MISSING_FROM_CRITICAL_STRINGS — <b>dropped</b> | SCOPED — <b>exports</b> |
| 04 | SCOPE_MISSING_FROM_QUESTION — <b>dropped</b>         | SCOPED — <b>exports</b> |
| 05 | SCOPE_MISSING_FROM_QUESTION — <b>dropped</b>         | SCOPED — <b>exports</b> |

The second condition is the one with teeth. A qualifier that appears in the question but in no critical string is decoration: nothing downstream reads it. Requiring it in the critical strings is what turns this from a style rule into a gate.

A FINDING THE REVIEWER MUST DECIDE — G-STRICTER-THAN-BATCH-006

The preregistered rule requires the scope qualifier in the question AND in the critical strings. GOLD-B006-08 carries 'OpenAI Python SDK' in its question but not in its critical strings, and the owner approved it. Implemented as preregistered rather than loosened; recorded so the reviewer can decide whether batch 007 should hold to the stricter rule.

*recorded as a passing test, not tuned away.*

## 5. The calibration pilot was not run

**BLOCKED**

BLOCKED — the frozen evidence the pilot must draw from is not present in this environment

The preregistration fixes the pilot's input exactly: *10 evidence spans that failed batch 006 ONLY because no builder could express them — NO\_BUILDER / UNBUILDABLE*. That input cannot be obtained here, on four independent grounds.

- the NO\_BUILDER/UNBUILDABLE set was counted, never persisted.** batch 006 recorded removed.unbuildable = 2482 as an integer; scripts/export\_batch\_006.py:733 increments the counter and returns, discarding the fact's identity. No artifact records which spans they were.
- the corpus text exists only in Postgres, and this container has no corpus.** load\_docs() reads document\_version.normalized\_text joined to corpus\_snapshot\_version. The local cluster starts but holds only postgres/template0/template1 — there is no 'rag' database and no snapshot.
- the raw documents are not in the repository.** data/raw/ contains only .gitkeep and is gitignored by design; data/cache/ is empty.
- re-fetching would not reproduce the frozen snapshot.** data/manifests/v1-openai-anthropic.yaml lists 202 documents captured 2026-08-17 with no text and no content hashes. Re-fetching returns the documentation as it stands today, so offsets and evidence hashes would not match snap\_689e336380a054d8039dc35b2c09cd0a, and the evidence would not be the frozen evidence.

No pilot case was authored. Authoring 10 cases against invented or re-fetched evidence would produce exactly what the preregistration exists to prevent — a benchmark testing what its author imagined rather than what the documentation says — and the pilot's four thresholds would measure nothing.

### The four thresholds, unmeasured

| criterion                            | threshold  | measured                            |
|--------------------------------------|------------|-------------------------------------|
| independently judged factually sound | >= 8 of 10 | <b>not measured — pilot not run</b> |
| unsupported claims                   | 0          | <b>not measured — pilot not run</b> |
| relation direction reversals         | 0          | <b>not measured — pilot not run</b> |
| scope broadening                     | 0          | <b>not measured — pilot not run</b> |

**To unblock:** Restore the corpus snapshot into Postgres (or restore data/raw/ and re-ingest), then re-run batch 006's miners to re-derive the spans that reach no builder and take the pilot's 10 from that set.

Until the pilot runs and is independently reviewed, the paraphrasing lane does not scale. The preregistration is explicit that a failed pilot means revising the authoring contract and re-piloting — not proceeding — and an *absent* pilot is not a weaker condition than a failed one.

## 6. Invariants

---

- `retrieval_was_not_run` is still `true` and `systems_executed` is still `[]`. No retrieval system was run against any candidate at any point; SYSTEM-A and SYSTEM-B remain frozen and unexecuted.
- Closed batches modified: **0**. Dataset records modified: **0**. Eligibility state modified: **false**.
- Validation and holdout were neither inspected nor modified; no split is frozen.
- `human_verified` set by this work: **0**. Only the project owner. No AI may set `human_verified`, and controlled paraphrasing does not make Claude authoritative about anything.
- Files added, none modified:
  - `src/rag_v1/gold/factidentity.py` (new)
  - `src/rag_v1/gold/reasoningtype.py` (new)
  - `src/rag_v1/gold/questionscope.py` (new)
  - `tests/test_gold001_b007_fixes.py` (new)

**Next:** Independent review of the three fixes, and a decision on the G-STRICTER finding, before the corpus is restored and the pilot is run. The paraphrasing lane does not scale until the pilot passes.

---

Generated by `scripts/build_batch_007_efg_pdf.py` from the E/F/G implementation record, the batch-007 preregistration, the project-wide eligibility status and the closed batch-005 and batch-006 records. Every figure is read from those artifacts at build time. The build refuses to run if the record's state disagrees with the eligibility status, if the E/F/G checks re-run here disagree with the recorded results, if a batch-007 candidate or pilot artifact exists, or if the page would claim the pilot ran. Raw provider documentation is not redistributed; quoted spans are the short excerpts under review.