

NATQ-002 — benchmark identity and Stage-3 status

Rename complete and byte-verified · Stage 3 at 8 of 150 · EXP-029A cannot execute · branch head `3e9fbfc` · 2026-09-04

Three findings, three different confidence levels. The rename is **done and proven** — every content file is byte-identical before and after. Stage 3 is **in progress** at 8/150 and the pace is the binding constraint. EXP-029A **cannot execute**: its measurement artifact does not exist, so no result is available from anyone. Nothing in this document is a measurement of NATQ-002 — it has not been scored by any system.

8/150
STAGE-3 ADJUDICATED

6
SUPPORTED

2
REJECTED ON PREMISE

0
SYSTEMS RUN VS NATQ-002

1 — Benchmark identity

The identifier `NATQ-001` already belongs to the frozen benchmark. Two distinct objects cannot share it, so the benchmark under construction here is now **NATQ-002**. They are **not merged**.

	NATQ-001 (frozen, upstream)	NATQ-002 (under construction, here)
Cases	100 accepted	150 authored → target 100
Validation	40 — EXPOSED to development and model selection	not yet split
Holdout	60 — PRISTINE, unopened	not yet split
Case-ID shape	<code>NATQ-C-nnn</code>	<code>A01...E30</code>
Scored by any system	yes — validation slice	no

No case-ID collision. The namespaces are disjoint, so no ID was renamed and no migration map was required. The qualified citation form `NATQ2-<case>`, e.g. `NATQ2-A01` is recorded for cross-benchmark tables; the stored `case` field is deliberately unchanged, because rewriting adjudicated packets to fix a citation convention would be a content revision.

2 — The rename is provenance-only, and that is verifiable

Change class `PROVENANCE_ONLY`. Benchmark content revision: **False**. Question text altered: **False**. Adjudicated evidence altered: **False**.

Content files — byte-identical before and after

File	SHA256	State
authoring/raw-slice-A.json	299351e39373f35b...	identical
authoring/raw-slice-B.json	6f765a9b7d56a21e...	identical
authoring/raw-slice-C.json	16a34ac8f065fc07...	identical
authoring/raw-slice-D.json	1618da10d2b3d223...	identical
authoring/raw-slice-E.json	f1f06a56e19f4b71...	identical
authoring/verification-pass1.json	1391de3bd89b4fdc...	identical
stage3/identifier-sweep.json	1971ccfb32f43569...	identical
stage3/packets-batch-01-partial.json	3a8a2934d5619ba0...	identical

Identity files — the only two that changed

File	Superseded	Current
NATQ-002-manifest.json	b167854268b334d3...	c86080e086a5bcfd...
NATQ-002-protocol.md	103b81e05b8368bb...	deffee18d9d40da5...

Historical records were left alone. experiments/STATUS-V2-001.json, experiments/STATUS-V2-001.pdf, scripts/build_status_v2_pdf.py still carry their NATQ-001 wording. They are dated, circulated documents describing this benchmark under its former name; rewriting them would falsify what was reported at the time. The rename provenance record is the pointer that disambiguates them.

3 — Stage 3: 8 of 150 adjudicated

Case	Question	Status	Evidence shape	Spans
A09	does the model ever call two tools in one response	SUPPORTED	single_span_api_reference	1
A04	what happens if my json schema has additionalProperties true w	SUPPORTED	single_span_constraint_list	1
A05	screenshot resolution for computer use — does it matter what s	SUPPORTED	single_span_procedure	1
C12	max_tokens deprecated in favor of max_completion_tokens?	REJECT	None	0
A02	tool_choice required vs any — whats the difference	SUPPORTED	multi_span_cross_provider	2
A03	I want the model to return ONLY valid json no preamble text	SUPPORTED	multi_span_procedure	2
A08	getting an error that my schema root must be an object	REJECT	—	0
A10	how do i turn off parallel tool calls	SUPPORTED	single_span_api_reference	1

A02 forced a validator change. No single document defines both `required` (OpenAI Agents SDK) and `any` (Anthropic Messages) — they are different vendors' names for forcing a tool call. The validator assumed one `version_id` per packet, so spans now carry their own and a packet whose evidence crosses providers must declare `cross-provider` and list them, rather than claim one provider while citing another. Batch 01 re-validates unchanged.

Two rejections, both on undocumented premises. A08 asks about a “schema root must be an object” error; the corpus enumerates every schema constraint and states no root-type rule. The phrase `must be an object` does occur — but only about input *messages*, a different subject, so anchoring there would have been a false positive that taught the benchmark to reward a wrong retrieval.

Two risks the coordinator should price in now.

Pace. 142 cases remain. The token-coverage locator ranks poorly on conversational questions, so most cases need literal probes and direct document reads — roughly two to four tool calls each. This is several sessions of work, not one.

Yield. 2 of 8 adjudicated cases were rejected because the corpus never states their premise. If that rate holds, 150 authored questions will not yield 100 accepted cases, and topping up means another cold-context authoring round. A rate on the full slice A will follow.

4 — Duplicates must be detected by evidence, not wording

A lexical screen over all 150 questions at Jaccard ≥ 0.34 found 2 pairs — and **missed the one that matters**.

Case	Question	Gold span
A09	does the model ever call two tools in one response	ver_18c692f4d28bd01c0a5cac553fcf01a7 @35392-35574
A10	how do i turn off parallel tool calls	ver_18c692f4d28bd01c0a5cac553fcf01a7 @35392-35574

The two share almost no vocabulary yet resolve to a **byte-identical span**. Two accepted cases resolving to one span double that span's weight in every aggregate metric, and a system that happens to retrieve it well gets paid twice for one capability.

Rule adopted. after all 150 are adjudicated, group SUPPORTED cases by their gold span set; any group of size > 1 is an evidence collision. Each collision gets an explicit ruling before the accepted 100 are frozen. Critically: **do not resolve collisions by looking at how any retrieval system scores the cases**. Evidence-collision detection cannot run before adjudication, because the gold span being compared does not exist until a case is adjudicated — so dedup is a gate *between* adjudication and selection, not a pre-screen. Both A09 and A10 are individually valid, recorded PASS, and flagged; nothing is deleted.

5 — EXP-029A cannot execute

Status `CANNOT_EXECUTE` . **Scores emitted:** none, and none estimated. The blocker is not a missing file this session failed to find — the operator confirmed `EXP-029A-raw-outputs.jsonl` exists nowhere, so the generation half has never been run by anyone.

This container cannot run it either: `cuda_available` False, `torch_installed` False, `transformers_installed` False, huggingface.co unreachable (`CONNECT` tunnel failed, response 403). The same two blockers stopped EXP-025A.

The analysis half is ready and unblocked

- All 23 STOP-condition checks from the EXP-029A specification.
- The ten preregistered gate conditions.
- The locked-16 recovery table against the NATQ-DIAG-002 locked span set.
- Preservation check against the EXP-026A all-pointwise oracle recognized-span set.
- The canonical multi-span diagnosis under METRIC-DEFS-V2 (12 cases / 25 spans).

- Micro-MRR over the canonical 53 gold spans.

Unblock condition. A completed EXP-029A generation run on a CUDA host producing EXP-029A-raw-outputs.jsonl, delivered with its SHA256. Until that file exists and its hash is recorded, EXP-029A has no result and no SUPPORTED value.

6 — Relayed lineage: indexed, not verified

The upstream artifacts pasted into chat are recorded as a receipt with `verified_here=False` and `byte_verified=False` . Their bytes were **not** reconstructed into this repository — doing so would create a second authority for artifacts this session never saw computed.

What is verified here:

- METRIC-DEFS-V2 canonical definitions match the EXP-024A Phase 0 reconciliation performed independently in this session.
- EXP-026A cites artifacts produced and hashed in this session as its inputs.
- METRIC-DEFS-V2 credits this session's EXP-023A artifacts.

Closed-state chain as relayed (directions unambiguous; counts transcribed, not verified): EXP-023A SUPPORTED=false → EXP-024A SUPPORTED=false → EXP-025A SUPPORTED=false → EXP-026A SUPPORTED=false → NATQ-DIAG-002 diagnostic, no SUPPORTED value → EXP-027A SUPPORTED=false → EXP-028A SUPPORTED=false → EXP-029A no result — measurement artifact does not exist.

If EXP-029A analysis is assigned here, those artifacts must be pushed to a branch, not pasted. The analysis compares against locked span sets and hash-identified populations; working from retyped numbers would fail the provenance bar this program has held everywhere else.

7 — Constraints honoured

Constraint	State
natq002 scored	False
holdout opened	False
split decided	False
systems overwritten	False
frozen checkout edited	False
scores invented	False

Still gated, in order: finish authoring and adjudication; run the dedup gate; freeze the accepted 100; hash after identity normalisation; *then* return to the coordinator for the split and preregistration. The split must not be chosen by looking at system results, and no retrieval or reranking architecture has been run against NATQ-002.