

NATQ-002 — Stage-3 adjudication report

22 of 150 adjudicated · 16 supported, 6 rejected, 1 deferred · 16 validator PASS / 0 FAIL · head `c5833dc` · 2026-09-04

Read the coverage before the rate. The rejection rate below is 6/22 (27.3%), but slice A supplies 22 of the 23 cases touched so far. The rate is currently a property of *one authoring slice*, not of the benchmark. Slice A skews toward troubleshooting scenarios that reference documentation does not cover, which is exactly the kind of question that rejects. **No yield decision should be taken on this sample.**

22/150
ADJUDICATED

16
SUPPORTED

6
REJECTED ON
PREMISE

0
VALIDATOR
FAILURES

0
SYSTEMS RUN

1 — Coverage by authoring slice

Slice	Done	Total	Progress
A	22	30	
B	0	30	
C	1	30	
D	0	30	
E	0	30	

This table is the reason the yield number is not yet actionable. Slices B, D and E are entirely untouched and slice C has one case. Each slice was authored by a different isolated cold-context agent from a different domain prompt, so their premise-support rates are not interchangeable.

2 — Supported cases (16)

Case	Question	Document	Provider	Spans
A09	does the model ever call two tools in one response	Messages	anthropic	1
A04	what happens if my json schema has additionalProperties tr	Structured outputs	anthropic	1
A05	screenshot resolution for computer use — does it matter wh	Coordinates and bounding b	anthropic	1
A02	tool_choice required vs any — whats the difference	Messages + Agents (cross-d	cross-provider	2
A03	I want the model to return ONLY valid json no preamble tex	Structured outputs	anthropic	2
A10	how do i turn off parallel tool calls	Messages	anthropic	1
A07	agents sdk handoff vs just giving it a tool that calls the	Tools	openai	2
A12	what do i send back after a tool_use block	Handle tool calls	anthropic	3
A13	tool result error — how to tell the model the tool failed	Handle tool calls	anthropic	2
A14	max token limit hit in the middle of a tool call, what do	Stop reasons and fallback	anthropic	2
A19	bash tool and text editor tool — do i have to implement th	Tool use with Claude	anthropic	2
A16	is there a way to make one specific tool always be the fir	Messages	anthropic	1
A21	streaming + tool calls = arguments come in fragments?	Fine-grained tool streamin	anthropic	1
A23	enum in schema not being respected	Structured outputs	anthropic	1
A25	recursive schema (tree of nodes) supported?	Structured outputs	anthropic	2
A27	prefilling assistant turn with { to force json — is that s	Prompting best practices	anthropic	1

Cross-provider: A02 — evidence spans two providers' documents, which the validator now expresses through span-level `version_id` plus an explicit `cross-provider` declaration.

3 — Rejected on premise (6)

These are not failures to find evidence. They are questions whose premise the frozen corpus never states, so no gold span can exist. Admitting them would train the benchmark to reward a system for retrieving a passage that does not answer the question.

Case	Question	Why rejected
C12	max_tokens deprecated in favor of max_completion_tok	UNSUPPORTED - the identifier max_completion_tokens does not occur anywhere in the frozen corpus (0 occurrences across
A08	getting an error that my schema root must be an obje	Premise not documented in the frozen corpus.
A01	how do I stop claude from calling the same tool over	Premise not documented in the frozen corpus.
A06	migrating from functions/function_call to tools/tool	Premise not documented in the frozen corpus.
A15	CU model clicking wrong coordinates on retina displa	Premise not documented in the frozen corpus.
A18	what's the max number of tools i can define	Premise not documented in the frozen corpus.

The near-miss discipline. Each rejection records the probes run and, where adjacent text exists, why anchoring there would be a false positive. A01 is the clearest case: `loop` appears 41 times across two documents, but always meaning the agentic request/response cycle — never a model repeating a call. A15 is flagged in its own packet as the rejection most open to being overturned, since the chain retina → high resolution → heavy downscaling → imprecise clicks has every link documented except the first.

4 — Deferred, not forced (1)

A11 — claude keeps saying "I'll use the tool" in text before actually calling it, can I suppress that

The corpus separately (a) gives general guidance to instruct the model to respond without preamble, and (b) shows preamble text preceding a `tool_use` block. It never connects them, and `preamble` is absent from every tool-use document checked.

Admitting it means citing general preamble guidance for a tool-use question the corpus never scopes it to. Rejecting it means calling a question unsupported when a directly applicable instruction does exist. Either verdict is a judgement the coordinator should make, not one to bury in a packet. It is counted in neither the supported nor the rejected column.

5 — Evidence collisions

Cases	Shared gold span
A09, A10	ver_18c692f4d28bd01c0a... @35392-35574

Detected by gold span, not by wording: a lexical screen over all 150 questions missed this pair entirely because the two share almost no vocabulary. The full collision gate runs **after all 150 are adjudicated, group SUPPORTED cases by their gold span set; any group of size > 1 is an evidence collision, and do not resolve collisions by looking at how any retrieval system scores the cases.**

6 — Yield: projection shown, and why it is not yet a forecast

Quantity	Value
Observed rejection rate	6/22 (27.3%)
Naive projection of accepted cases at that rate	109
Target	100
Sample is representative	False

109 clears 100, but only by 9 — and that margin is consumed before the collision gate removes any duplicate-evidence cases. Deferred cases also do not count toward the accepted pool unless the coordinator admits them. The honest reading is that the benchmark is **near the line, on a biased sample**, and the real number arrives when slice A closes and slices B–E begin.

7 — Gates still closed

Gate	State
all 150 complete	False
final collision gate run	False
hundred frozen	False
benchmark hashed	False
split decided	False

Constraints held

Constraint	State
natq002 scored	False
holdout opened	False
split decided	False
question wording altered	False
systems run	False

Order of remaining work: finish the 127 untouched questions; run the collision gate; freeze the accepted 100; hash after identity normalisation; then return to the coordinator for the split. No retrieval or reranking architecture has been run against NATQ-002, and the NATQ-001 holdout has not been opened.